


Virtual Threads vs Platform Threads: Latency and Queue Length

(a) Mean Latency

PT $\rho=0.75$: 753 ms
(± 1417 ms) — off-scale



(b) Queue Length

PT $\rho=0.75$: 232 req
(± 471) — off-scale

